

# Distributed Databases

CO527 Advanced Database Systems

1

## Overview

- Distributed Database Concepts
  - Advantages of Distributed Databases
- Data Fragmentation, Replication, and Allocation Techniques
- Concurrency Control and Recovery
- Distributed Transactions
  - Two-phase commit protocol
  - Three-phase commit protocol
- Distributed Query Processing
- Types of Distributed Database Systems

2

## Distributed Database Concepts

- A Distributed Database (DDB) is a collection of multiple logically related database distributed over a computer network, and a distributed database management system as a software system that manages a distributed database while making the distribution transparent to the user.
- Distributed Database is a combined result of the two technologies
  - Database technology
  - Network and Data Communication Technology

3


## Distributed Database Concepts (cont.)

- Recently distributed and database technologies are being developed for dealing with vast amount of data which is known as big data technologies.
- Then data mining and machine learning algorithms are used to extract needed knowledge.

4

## Advantages of Distributed Database System


- Management of distributed data with different **levels of transparency**:
  - This refers to the physical placement of data (files, relations, etc.) which is not known to the user (distribution transparency/data organization transparency).



5

## Advantages of Distributed Database System

- Example
  - The EMPLOYEE, PROJECT, and WORKS\_ON tables may be fragmented horizontally and stored with possible replication as shown below.



6

## Advantages of Distributed Database System

- **Distribution and Network transparency:**
  - Users do not have to worry about operational details of the network.
  - There is Location transparency, which refers to freedom of issuing command from any location without affecting its working.
  - Then there is Naming transparency, which allows access to any names object (files, relations, etc.) from any location.
- **Replication transparency:**
  - It allows to store copies of a data at multiple sites as shown in the above diagram.
  - This is done to minimize access time to the required data.
- **Fragmentation transparency:**
  - Allows to fragment a relation horizontally (create a subset of tuples of a relation) or vertically (create a subset of columns of a relation).

7

## Advantages of Distributed Database System

- Other Advantages
  - **Improved ease and flexibility of application development**
  - **Increased reliability and availability:**
    - Reliability refers to system live time, that is, system is running efficiently most of the time. Availability is the probability that the system is continuously available (usable or accessible) during a time interval.
    - A distributed database system has multiple nodes (computers) and if one fails then others are available to do the job.
  - **Improved performance:**
    - A distributed DBMS fragments the database to keep data closer to where it is needed most.
    - This reduces data management (access and modification) time significantly.
  - **Easier expansion (scalability):**
    - Horizontal scalability: Allows new nodes (computers) to be added anytime without chaining the entire configuration.
    - Vertical scalability: Expanding capacity of individual nodes of the system or processing power.

8

## Data Fragmentation, Replication and Allocation

- **Data Fragmentation**
  - Split a relation into logically related and correct parts. A relation can be fragmented in two ways:
    - Horizontal Fragmentation
    - Vertical Fragmentation

9

## Data Fragmentation, Replication and Allocation

- **Horizontal fragmentation (Sharding)**
  - It is a horizontal subset of a relation which contain those of tuples which satisfy selection conditions.
  - Consider the Employee relation with selection condition (DNO = 5). All tuples satisfy this condition will create a subset which will be a horizontal fragment of Employee relation.
  - A selection condition may be composed of several conditions connected by AND or OR.
  - Derived horizontal fragmentation: It is the partitioning of a primary relation to other secondary relations which are related with Foreign keys.

10

## Data Fragmentation, Replication and Allocation

- **Vertical fragmentation**
  - It is a subset of a relation which is created by a subset of columns. Thus a vertical fragment of a relation will contain values of selected columns. There is no selection condition used in vertical fragmentation.
  - Consider the Employee relation. A vertical fragment of can be created by keeping the values of Name, Bdate, Sex, and Address.
  - Because there is no condition for creating a vertical fragment, each fragment must include the primary key attribute of the parent relation Employee. In this way all vertical fragments of a relation are connected.

11

## Data Fragmentation, Replication and Allocation

- **Data Replication**
  - Database is replicated to all sites.
  - In full replication the entire database is replicated (called fully replicated distributed database) and in partial replication some selected part is replicated to some of the sites.
- **Data Distribution (Data Allocation)**
  - This is relevant only in the case of partial replication or partition.
  - The selected portion of the database is distributed to the database sites.

12

## Data Fragmentation, Replication and Allocation

(a) EMPD, 5

Empno	Mval	Lname	Sal	Salary	Super_ssn	Dep
John	B	Smith	123456789	30000	333445555	5
Franklin	T	Wong	333445555	40000	888665555	5
Ramesh	K	Narayan	666884444	38000	333445555	5
Joyce	A	English	453453453	25000	333445555	5

Figure 25.8 Allocation of fragments to sites. (a) Relation fragments at site 2 corresponding to department 5. (b) Relation fragments at site 3 corresponding to department 4.

DEP 5

Empno	Department	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22

DEP 5 LOCs

Department	Location
5	Bellare
5	Sugarland
5	Houston

WORKS ON 5

Empno	Proj	Hours
123456789	1	52.5
123456789	2	75
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0

PROJ 5

Project	Prj_start	Prj_end	Prj_mgr
Product X	1	Endless	5
Product Y	3	Sugarland	5
Product Z	3	Houston	5

Data at site 2

13

## Data Fragmentation, Replication and Allocation

(b) EMPD, 4

Empno	Mval	Lname	Sal	Salary	Super_ssn	Dep
Alice	J	Zeloya	999887777	25000	987654321	4
Jennifer	S	Wolfe	987654321	43000	888665555	4
Ahmed	V	Jalilov	987654321	25000	987654321	4

DEP 4

Empno	Department	Mgr_ssn	Mgr_start_date
Administration	4	987654321	1995-01-01

DEP 4 LOCs

Department	Location
4	Stafford

WORKS ON 4

Empno	Proj	Hours
333445555	10	10.0
999887777	30	30.0
999887777	10	10.0
987654321	10	25.0
987654321	30	5.0
987654321	30	20.0
987654321	20	15.0

PROJ 4

Project	Prj_start	Prj_end	Prj_mgr
Computerization	10	Stafford	4
New benefits	30	Stafford	4

Data at site 3

14


## Concurrency Control and Recovery

- Dealing with multiple copies of data items:
  - The concurrency control must maintain global consistency. Likewise the recovery mechanism must recover all copies and maintain consistency after recovery.
- Failure of individual sites:
  - Database availability must not be affected due to the failure of one or two sites and the recovery scheme must recover them before they are available for use.
- Communication link failure:
  - This failure may create network partition which would affect database availability even though all database sites may be running.
- Distributed commit:
  - A transaction may be fragmented and they may be executed by a number of sites. This requires a two or three-phase commit approach for transaction commit.
- Distributed deadlock:
  - Since transactions are processed at multiple sites, two or more sites may get involved in deadlock. This must be resolved in a distributed manner.

16

## Concurrency Control and Recovery

- Distributed Concurrency control based on a distributed copy of a data item
  - Primary site technique:** A single site is designated as a primary site which serves as a coordinator for transaction management.



17

## Concurrency Control and Recovery

- Transaction management:
  - Concurrency control and commit are managed by this site.
  - In two phase locking, this site manages locking and releasing data items. If all transactions follow two-phase policy at all sites, then serializability is guaranteed.

18

## Concurrency Control and Recovery

- Transaction Management
  - Advantages:
    - An extension to the centralized two phase locking so implementation and management is simple.
    - Data items are locked only at one site but they can be accessed at any site.
  - Disadvantages:
    - All transaction management activities go to primary site which is likely to overload the site.
    - If the primary site fails, the entire system is inaccessible.
  - To aid recovery a backup site is designated which behaves as a shadow of primary site. In case of primary site failure, backup site can act as primary site.

19

## Concurrency Control and Recovery

- **Primary Copy Technique:**
  - In this approach, instead of a site, a data item partition is designated as primary copy. To lock a data item just the primary copy of the data item is locked.
- **Advantages:**
  - Since primary copies are distributed at various sites, a single site is not overloaded with locking and unlocking requests.
- **Disadvantages:**
  - Identification of a primary copy is complex. A distributed directory must be maintained, possibly at all sites.

20

## Concurrency Control and Recovery

- **Recovery from a coordinator failure**
  - In both approaches a coordinator site or copy may become unavailable. This will require the selection of a new coordinator.
- **Primary site approach with no backup site:**
  - Aborts and restarts all active transactions at all sites. Elects a new coordinator and initiates transaction processing.
- **Primary site approach with backup site:**
  - Suspends all active transactions, designates the backup site as the primary site and identifies a new back up site. Primary site receives all transaction management information to resume processing.
- **Primary and backup sites fail or no backup site:**
  - Use election process to select a new coordinator site.

21

## Concurrency Control and Recovery

- **Concurrency control based on voting:**
  - There is no primary copy of coordinator.
  - Send lock request to sites that have data item.
  - If majority of sites grant lock then the requesting transaction gets the data item.
  - Locking information (grant or denied) is sent to all these sites.
  - To avoid unacceptably long wait, a time-out period is defined. If the requesting transaction does not get any vote information then the transaction is aborted.

22

## Distributed Transactions

- Transaction that updates data on two or more systems
- **Challenge**
  - Handle machine, software, & network failures while preserving transaction integrity

23

## Distributed Transactions

- Each computer runs a transaction manager
  - Responsible for subtransactions on that system
  - Transaction managers communicate with other transaction managers
  - Performs prepare, commit, and abort calls for subtransactions
- Every subtransaction must agree to commit changes before the transaction can complete


24

## Transactional Commits are Consensus

- **Remember consensus?**
  - Agree on a value proposed by at least one process
- **The coordinator proposes to commit a transaction**
  - Participants will agree  $\Rightarrow$  all participants then commit
  - Participants will not agree  $\Rightarrow$  all participants then abort
- Here, we need unanimous agreement to commit

25

## Two-Phase Commit Protocol




26

## Two-phase commit protocol

- Goal:
  - *Reliably agree to commit or abort a collection of sub-transactions*
- All processes in the transaction will agree to commit or abort
- One transaction manager is elected as a **coordinator** – the rest are **participants**
- Assume:
  - **write-ahead log** in stable storage
  - No system dies forever
  - Systems can always communicate with each other

27

## Transaction States




When a participant enters the *prepared* state, it contacts the coordinator to start the commit protocol to commit the entire transaction

28

## Two-Phase Commit Protocol


- Phase 1: Voting Phase
  - Get commit agreement from *every participant*



29

## Two-Phase Commit Protocol

- Phase 1: Voting Phase
  - Get commit agreement from *every participant*




A single "no" response means that we will have to abort the transaction

30

## Two-Phase Commit Protocol

- Phase 2: Commit Phase
  - Send the results of the vote to every participant




Send abort if any participant voted "no"

31

## Two-Phase Commit Protocol

- Phase 2: Commit Phase
  - Get "I have committed" acknowledgements from *every participant*



32

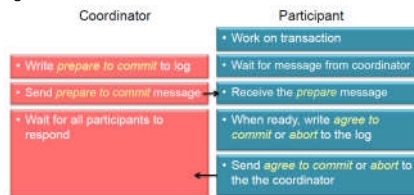
## Dealing with failure

- 2PC assumes a fail-recover model
  - Any failed system will eventually recover
- A recovered system cannot change its mind
  - If a node agreed to commit and then crashed, it must be willing and able to commit upon recovery
- Each system will use a writeahead (transaction) log
  - Keep track of where it is in the protocol (and what it agreed to)
  - As well as values to enable commit or abort (rollback)

33

## Two-Phase Commit Protocol: Phase 1

### Voting Phase

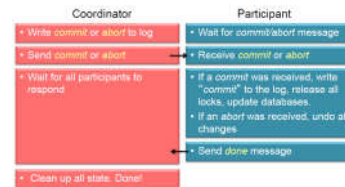


- Get distributed agreement: the coordinator asked each participant if it will commit or abort and received replies from each participant.

34

## Two-Phase Commit Protocol: Phase 2

### Commit Phase



- Tell *all participants* to *commit* or *abort*
- Get everyone's response that they're done

35

## Dealing with failure

### Failure during Phase 1 (voting)

- Coordinator dies
  - Some participants may have responded; others have no clue
    - Coordinator restarts; checks log; sees that voting was in progress
    - Coordinator *restarts voting*
- Participant dies
  - The participant may have died before or after sending its vote to the coordinator
    - If the coordinator received the vote, wait for other votes and go to phase 2
    - Otherwise: *wait for the participant* to recover and respond (keep querying it)

36

## Dealing with failure

### Failure during Phase 2 (commit/abort)

- Coordinator dies
  - Some participants may have been given commit/abort instructions
    - Coordinator *restarts*; checks log; *informs all participants* of chosen action
- Participant dies
  - The participant may have died before or after getting the commit/abort request
    - Coordinator *keeps trying to contact the participant* with the request
    - Coordinator *recovers*; checks log; gets request from coordinator
      - If it committed/aborted, acknowledge the request
      - Otherwise, process the commit/abort request and send back the acknowledgement

37

## Adding a recovery coordinator

- Another system can take over for the coordinator
  - Could be a participant that detected a timeout to the coordinator
- Recovery node needs to find the state of the protocol
  - Contact ALL participants to see how they voted
  - If we get **voting results from all participants**
    - We know that Phase 1 has completed
    - If all participants voted to commit  $\Rightarrow$  send commit request
    - Otherwise send abort request
  - **If ANY participant has not voted**
    - We know that Phase 1 has *not* completed
      - Restart the protocol
- But ... if a participant node also crashes, we're stuck!
  - Have to wait for recovery

38

## What's wrong with the 2PC protocol?

- Biggest problem: **it's a blocking protocol**
  - If the coordinator crashes, participants have no idea whether to commit or abort
    - A recovery coordinator helps in some cases
  - A non-responding participant will also result in blocking
- When a participant gets a commit/abort message, it does not know if every other participant was informed of the result

39

## Three-Phase Commit Protocol

- Same setup as the two-phase commit protocol:
  - Coordinator & Participants
- Add timeouts to each phase that result in an abort
- Propagate the result of the commit/abort vote to each participant *before telling them to act on it*
  - This will allow us to recover the state if any participant dies

40

## Three-Phase Commit Protocol

- Split the second phase of 2PC into two parts:
  - **2a. "Precommit" (prepare to commit) phase**
    - Send **Prepare** message to all participants when it received a yes from all participants in phase 1
    - Participants can prepare to commit but cannot do anything that cannot be undone
    - Participants reply with an acknowledgement
    - Purpose: let every participant know the state of the result of the vote so that state can be recovered if anyone dies
  - **2b. "Commit" phase (same as in 2PC)**
    - If coordinator gets ACKs for all "precommit" messages
      - It will send a **commit** message
    - Else it will abort – send an **abort** message to all participants

41

## Three-Phase Commit Protocol

- **Phase 1: voting phase**
  - Coordinator sends **canCommit?** queries to participants & et responses
  - Purpose: Find out if everyone agrees to commit
  - If the coordinator gets a **timeout** from any participant, or any **NO** replies are received
    - Send an **abort** to all participants
  - If a participant times out waiting for a request from the coordinator
    - It **aborts** itself (assume coordinator crashed)
  - Else continue to phase 2
- **Phase 2: Prepare to commit phase**
  - Send a **Prepare** message to all participants. Get **OK** messages from all participants
  - Purpose: let all participants know the decision to commit
  - If coordinator times out: assume participant crashed, send **abort** to all participants
    - The coordinator cannot count on every participant having received the Prepare message
- **Phase 3: finalize**
  - Send **commit** messages to participants and get responses from all
  - If participant times out: contact any other participant and move to that state (commit or abort)
  - If coordinator times out: that's ok

42

## 3PC Recovery

- If the coordinator crashes
  - A recovery node can query the state from any available participants
- Possible states that the participant may report:
  - **Already committed**
    - That means that every other participant has received a Prepare to Commit
    - Some participants may have committed
    - Send **Commit** message to all participants (just in case they didn't get it)
  - **Not committed but received a Prepare message**
    - That means that all participants agreed to commit; some may have committed
    - Send **Prepare to Commit** message to all participants (just in case they didn't get it)
    - Wait for everyone to acknowledge; then **commit**
  - **Not yet received a Prepare message**
    - This means no participant has committed; some may have agreed
    - Transaction can be **aborted** or the commit protocol can be **restarted**

43

## 3PC Weaknesses

- Main weakness of 3PC
  - May have problems when the network gets partitioned
  - Partition A: nodes that received *Prepare message*
    - Recovery coordinator for A: allows commit
  - Partition B: nodes that did not receive *Prepare message*
    - Recovery coordinator for B: aborts
  - Either of these actions are legitimate as a whole
    - But when the network merges back, the system is inconsistent
- Not good when a crashed coordinator recovers
  - It needs to find out that someone took over and stay quiet
  - Otherwise it will mess up the protocol, leading to an inconsistent state

44

## 3PC coordinator recovery problem

- Suppose a coordinator sent a Prepare messages to all participants
- Suppose all participants acknowledged these message
  - BUT the coordinator died before it got all acknowledgements
- A recovery coordinator queries a participant
  - Continues with the commit: Sends Prepare, gets ACKs, sends Commit
- At the same time...
  - The original coordinator recovers
  - Realizes it is still missing some replies from the Prepare
  - Times out and decides to send an Abort to all participants
- Some processes may commit while others abort!
- 3PC is not resilient against asynchronous fail-recover failures

45

## Query Processing in Distributed Databases

### • Issues

- Cost of transferring data (files and results) over the network.
  - This cost is usually high so some optimization is necessary.
- Example relations: Employee at site 1 and Department at Site 2
  - Employee at site 1. 10,000 rows. Row size = 100 bytes. Table size =  $10^6$  bytes.

Fname	Minit	Lname	SSN	Bdate	Address	Sex	Salary	Superssn	Dno
-------	-------	-------	-----	-------	---------	-----	--------	----------	-----

- Department at Site 2. 100 rows. Row size = 35 bytes. Table size = 3,500 bytes.

Dname	Dnumber	Mgrssn	Mgrstartdate
-------	---------	--------	--------------

- Q: For each employee, retrieve employee name and department name Where the employee works.
- Q:  $\Pi_{Fname, Lname, Dname} (Employee \bowtie_{Dno = Dnumber} Department)$

46

## Query Processing in Distributed Databases

### • Result

- The result of this query will have 10,000 tuples, assuming that every employee is related to a department.
- Suppose each result tuple is 40 bytes long. The query is submitted at site 3 and the result is sent to this site.
- Problem: Employee and Department relations are not present at site 3.

Site 1:  
EMPLOYEE

Fname	Minit	Lname	Sex	Bdate	Address	Sex	Salary	Superssn	Dno
-------	-------	-------	-----	-------	---------	-----	--------	----------	-----

10,000 records  
each record is 100 bytes long  
Row field is 8 bytes long  
Dno field is 4 bytes long

Site 2:  
DEPARTMENT

Dname	Dnumber	Mgrssn	Mgr_start_date
-------	---------	--------	----------------

100 records  
each record is 35 bytes long  
Dnumber field is 4 bytes long  
Mgrssn field is 8 bytes long

47

## Query Processing in Distributed Databases

### • Strategies:

1. Transfer Employee and Department to site 3.
  - Total transfer bytes =  $1,000,000 + 3500 = 1,003,500$  bytes.
2. Transfer Employee to site 2, execute join at site 2 and send the result to site 3.
  - Query result size =  $40 * 10,000 = 400,000$  bytes. Total transfer size =  $400,000 + 1,000,000 = 1,400,000$  bytes.
3. Transfer Department relation to site 1, execute the join at site 1, and send the result to site 3.
  - Total bytes transferred =  $400,000 + 3500 = 403,500$  bytes.

- Optimization criteria: minimizing data transfer.
  - Preferred approach: strategy 3.

48

## Query Processing in Distributed Databases

### • Consider the query

- Q: For each department, retrieve the department name and the name of the department manager

### • Relational Algebra expression:

- $\Pi_{Fname, Lname, Dname} (Employee \bowtie_{Mgrssn = SSN} Department)$

49



## Query Processing in Distributed Databases

- The result of this query will have 100 tuples, assuming that every department has a manager, the execution strategies are:
  - Transfer Employee and Department to the result site and perform the join at site 3.
    - Total bytes transferred =  $1,000,000 + 3500 = 1,003,500$  bytes.
  - Transfer Employee to site 2, execute join at site 2 and send the result to site 3. Query result size =  $40 * 100 = 4000$  bytes.
    - Total transfer size =  $4000 + 1,000,000 = 1,004,000$  bytes.
  - Transfer Department relation to site 1, execute join at site 1 and send the result to site 3.
    - Total transfer size =  $4000 + 3500 = 7500$  bytes.
- Preferred strategy: Choose strategy 3.

50

## Query Processing in Distributed Databases

- Now suppose the result site is 2. Possible strategies :
  - Transfer Employee relation to site 2, execute the query and present the result to the user at site 2.
    - Total transfer size =  $1,000,000$  bytes for both queries Q and Q'.
  - Transfer Department relation to site 1, execute join at site 1 and send the result back to site 2.
    - Total transfer size for Q =  $400,000 + 3500 = 403,500$  bytes and for Q' =  $4000 + 3500 = 7500$  bytes.

51


## Query Processing in Distributed Databases

- Semijoin:
  - Objective is to reduce the number of tuples in a relation before transferring it to another site.
- Example execution of Q or Q':
  - Project the join attributes of Department at site 2, and transfer them to site 1. For Q,  $4 * 100 = 400$  bytes are transferred and for Q',  $9 * 100 = 900$  bytes are transferred.
  - Join the transferred file with the Employee relation at site 1, and transfer the required attributes from the resulting file to site 2. For Q,  $34 * 10,000 = 340,000$  bytes are transferred and for Q',  $39 * 100 = 3900$  bytes are transferred.
  - Execute the query by joining the transferred file with Department and present the result to the user at site 2.

52

## Types of Distributed Database Systems


- Homogeneous
  - All sites of the database system have identical setup, i.e., same database system software.
  - The underlying operating system may be different.
    - For example, all sites run Oracle or DB2, or Sybase or some other database system.
  - The underlying operating systems can be a mixture of Linux, Window, Unix, etc.



53

## Types of Distributed Database Systems

- Heterogeneous
  - Federated: Each site may run different database system but the data access is managed through a single conceptual schema.
    - This implies that the degree of local autonomy is minimum. Each site must adhere to a centralized access policy. There may be a global schema.
  - Multidatabase: There is no one conceptual global schema. For data access a schema is constructed dynamically as needed by the application software.



54

## Types of Distributed Database Systems

- Federated Database Management Systems Issues
  - Differences in data models:
    - Relational, Objected oriented, hierarchical, network, etc.
  - Differences in constraints:
    - Each site may have their own data accessing and processing constraints.
  - Differences in query language:
    - Some site may use SQL, some may use SQL-89, some may use SQL-92, and so on.

55